

Guide to Concurrency and Multithreading

1 Threads, Locks, and Deadlock

1.1 Threads

Threads are the smallest unit of execution within a process, allowing multiple tasks to run concurrently. Each thread shares the same memory space but has its own stack and program counter.

- **Creation:** In Java, extend `Thread` or implement `Runnable`. In Python, use `threading.Thread`.
- **Lifecycle:** Threads have states—new, runnable, blocked, waiting, timed waiting, and terminated.
- **Example (Python):**

```
1 import threading
2 def print_numbers():
3     for i in range(5):
4         print(f"Thread {threading.current_thread().name}: {i}")
5 thread = threading.Thread(target=print_numbers, name="NumberThread")
6 thread.start()
7 thread.join()
```

1.2 Locks

Locks ensure that only one thread accesses a critical section of code at a time, preventing race conditions.

- **Types:** Reentrant (allows the same thread to acquire the lock multiple times) and non-reentrant.
- **Usage:** In Java, use `synchronized` blocks or `ReentrantLock`. In Python, use `threading.Lock`.
- **Example (Java):**

```
1 import java.util.concurrent.locks.ReentrantLock;
2 class Counter {
3     private final ReentrantLock lock = new ReentrantLock();
4     private int count = 0;
5     public void increment() {
6         lock.lock();
7         try {
8             count++;
9         } finally {
10            lock.unlock();
11        }
12    }
13 }
```

1.3 Deadlock

Deadlock occurs when two or more threads are blocked forever, each waiting for a resource the other holds.

- **Conditions for Deadlock:**

- Mutual Exclusion: Resources are held exclusively.
- Hold and Wait: Threads hold at least one resource and wait for another.
- No Preemption: Resources cannot be forcibly taken.
- Circular Wait: Threads form a circular dependency.

- **Prevention:**

- Lock Ordering: Always acquire locks in a fixed order.
- Timeout: Use timeouts when acquiring locks.
- Avoid Nested Locks: Minimize lock dependencies.

- **Example (Deadlock in Java):**

```
1 class DeadlockExample {
2     static Object resource1 = new Object();
3     static Object resource2 = new Object();
4     public static void main(String[] args) {
5         Thread t1 = new Thread(() -> {
6             synchronized (resource1) {
7                 try { Thread.sleep(100); } catch (Exception e) {}
8                 synchronized (resource2) {
9                     System.out.println("Thread 1: Holding both
10                        resources");
11                 }
12             }
13         });
14         Thread t2 = new Thread(() -> {
15             synchronized (resource2) {
16                 try { Thread.sleep(100); } catch (Exception e) {}
17                 synchronized (resource1) {
18                     System.out.println("Thread 2: Holding both
19                        resources");
20                 }
21             }
22         });
23         t1.start();
24         t2.start();
25     }
26 }
```

2 Producer-Consumer Problems

The producer-consumer problem involves two types of threads: producers that generate data and consumers that process it, sharing a bounded buffer.

- **Key Components:**

- **Buffer:** A fixed-size queue (e.g., `ArrayBlockingQueue` in Java or `queue.Queue` in Python).
- **Producer:** Adds data to the buffer.
- **Consumer:** Removes and processes data from the buffer.
- **Challenges:**
 - Ensure producers do not add to a full buffer.
 - Ensure consumers do not remove from an empty buffer.
 - Prevent race conditions during access.
- **Solution:** Use synchronization mechanisms like locks, semaphores, or condition variables.
- **Example (Python):**

```

1 import threading
2 import queue
3 import time
4 q = queue.Queue(maxsize=10)
5 def producer():
6     for i in range(20):
7         q.put(i)
8         print(f"Produced: {i}")
9         time.sleep(0.1)
10 def consumer():
11     while True:
12         item = q.get()
13         print(f"Consumed: {item}")
14         q.task_done()
15         time.sleep(0.2)
16 producer_thread = threading.Thread(target=producer)
17 consumer_thread = threading.Thread(target=consumer, daemon=True)
18 producer_thread.start()
19 consumer_thread.start()
20 producer_thread.join()

```

3 Semaphores and Mutexes

3.1 Semaphores

Semaphores are synchronization primitives that control access to a resource by maintaining a count.

- **Binary Semaphore:** Acts like a mutex (0 or 1).
- **Counting Semaphore:** Allows a fixed number of threads to access a resource.
- **Usage:** In Java, use `Semaphore`. In Python, use `threading.Semaphore`.
- **Example (Python):**

```

1 import threading
2 semaphore = threading.Semaphore(2) % Allows 2 threads
3 def task(name):
4     with semaphore:
5         print(f"{name} acquired semaphore")
6         time.sleep(1)
7         print(f"{name} released semaphore")
8 threads = [threading.Thread(target=task, args=(f"Thread-{i}",)) for i
9             in range(5)]
10 for t in threads:
11     t.start()
12 for t in threads:
13     t.join()

```

3.2 Mutexes

A mutex (mutual exclusion) is a lock that ensures only one thread accesses a resource at a time.

- **Difference from Semaphore:** Mutex is typically used for mutual exclusion, while semaphores can control access for multiple threads.
- **Usage:** In C, use `pthread_mutex_t`. In Java, use `synchronized` or `ReentrantLock`.
- **Example (C):**

```

1 #include <pthread.h>
2 #include <stdio.h>
3 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
4 int count = 0;
5 void* increment(void* arg) {
6     pthread_mutex_lock(&mutex);
7     count++;
8     printf("Count: %d\n", count);
9     pthread_mutex_unlock(&mutex);
10    return NULL;
11 }
12 int main() {
13     pthread_t threads[5];
14     for (int i = 0; i < 5; i++)
15         pthread_create(&threads[i], NULL, increment, NULL);
16     for (int i = 0; i < 5; i++)
17         pthread_join(threads[i], NULL);
18     return 0;
19 }

```

4 Thread-Safe Collections

Thread-safe collections are designed to handle concurrent access without external synchronization.

- **Java Examples:**
 - `ConcurrentHashMap`: Allows concurrent reads and writes.
 - `CopyOnWriteArrayList`: Creates a new copy on write operations, suitable for read-heavy scenarios.

– ArrayBlockingQueue: Thread-safe bounded queue for producer-consumer.

- **Python Examples:**

– queue.Queue: Thread-safe queue for producer-consumer.

– collections.deque (with locks): Not inherently thread-safe but can be used with threading.Lock.

- **Example (Java):**

```
1 import java.util.concurrent.ConcurrentHashMap;
2 class ThreadSafeMap {
3     public static void main(String[] args) {
4         ConcurrentHashMap<String, Integer> map = new
5             ConcurrentHashMap<>();
6         Runnable task = () -> {
7             for (int i = 0; i < 100; i++) {
8                 map.compute("key", (k, v) -> v == null ? 1 : v + 1);
9             }
10            Thread t1 = new Thread(task);
11            Thread t2 = new Thread(task);
12            t1.start();
13            t2.start();
14            try {
15                t1.join();
16                t2.join();
17            } catch (InterruptedException e) {}
18            System.out.println("Final count: " + map.get("key"));
19        }
20    }
```